

PartnerScope Overview

PartnerScope

Intelligent Partner Discovery for Startups

The Problem

Finding the right business partners is one of the most time-consuming and critical challenges for startups. Traditional approaches involve:

- Hours of manual research across databases and websites
- Inconsistent evaluation criteria across team members
- No systematic way to compare candidates
- Difficulty articulating exactly what kind of partner is needed

PartnerScope solves this by combining conversational AI with structured search and evaluation—turning weeks of research into minutes.

How It Works

PartnerScope guides startups through a complete partner discovery workflow:

Discovery Chat → Partner Search → Evaluation → Comparison

1. Discovery Chat

Start with a conversation, not a form. Our AI coach helps you articulate your partnership needs through natural dialogue.

What makes it different: - Acts as a strategic advisor, not a data collector - Proactively suggests partnership types you may not have considered - Extracts structured requirements from conversation

Sample dialogue:

You: "We're building a wellness app for students"

PartnerScope: "Tell me more about how it works—what does the

app actually do, and what makes it different from other wellness apps?"

You: "It uses AI agents to create a personalized co-regulation companion that responds to check-ins and nudges users toward social contact"

PartnerScope: "That's a unique approach—AI-driven co-regulation rather than generic wellness content. For something this personalized, you'd likely need a Pilot Population Partner—an organization with users who could test it over several weeks."

What scale are you thinking for a pilot? 50 users? 500?"

What gets extracted: - Partner type (pilot, distribution, validation, etc.) - Minimum requirements ("must have 100+ participants") - Success criteria ("60% retention at week 8") - Red flags to avoid ("no single owner on the org side") - Information to collect during outreach

2. Partner Search

Multi-source search combines database lookups with real-time AI web search.

Data Sources:

Source	Best For	Speed
Database/CSV	Known players, established companies	Fast
AI Web Search	Emerging players, niche specialists	Thorough

Use both together for comprehensive coverage, or individually based on your needs.

5-Phase Search Architecture:

Phase	What Happens	Why It Matters
1. Discovery	4 search queries from different angles	Casts a wide net
2. Reflection	Analyzes gaps, searches for non-obvious partners	Finds what you didn't know to look for
3. Decomposition	Breaks need into sub-needs, targeted search for each	Ensures comprehensive coverage
4. Batch Filtering	Scores candidates in small batches	Maintains quality across large candidate sets
5. Enrichment	Fills missing data for top candidates	Clean, complete output

Output: 20 ranked partners with: - Company details (name, website, industry, location) - Partnership fit rationale - Needs they satisfy - Validation score (1-10)

3. Evaluation

Multi-dimensional assessment with human-in-the-loop refinement.

Dynamic Strategy: Instead of fixed criteria, you control the evaluation dimensions and weights:

Proposed Strategy:

1. Market Compatibility (25%)
2. Technical Synergy (25%)
3. Strategic Alignment (20%)
4. Growth Potential (15%)
5. Risk Profile (15%)

You: "Focus more on technical synergy, we need API integration"

Adjusted:

- Technical Synergy: 35% (was 25%)
- Market Compatibility: 20% (was 25%)

Available Dimensions:

Dimension	What It Measures
Market Compatibility	Target market alignment
Technical Synergy	Technology compatibility
Financial Health	Stability, funding status
Operational Capacity	Logistics, scale capability
Geographic Coverage	Regional presence
Strategic Alignment	Business goal fit
Cultural Fit	Organizational compatibility
Resource Complementarity	Complementary assets
Growth Potential	Scalability opportunity
Risk Profile	Potential challenges

Post-Evaluation Refinement: Adjust results without re-running the entire evaluation:

Action	Example	What Happens
Exclude	"Remove TechPartner, they're a competitor"	Filters out, re-ranks remaining
Reweight	"Prioritize geographic coverage"	Adjusts weights, recalculates scores
Filter	"Show only top 3"	Returns focused subset
Focus	"Tell me more about their weaknesses"	Deeper analysis

4. Comparison with External Research

Validate results against other AI research tools—Gemini Deep Research, OpenAI Deep Research, or Claude.

How it works: 1. Run your preferred external research tool with the same query 2. Paste the results into PartnerScope 3. PartnerScope evaluates them using the same criteria 4. Side-by-side comparison shows which found better partners

Metrics compared: - Top-8 average score - Overall average score - Score distribution - Specific candidate rankings

This provides objective validation that your partner search is producing high-quality results.

Quality vs. Speed vs. Cost

Choose the right balance for your situation:

Mode	Quality	Speed	Cost	Best For
Quality	Highest	~20-30 min	~\$0.32/search	Final selection, high-stakes decisions
Balanced	Good	~20 sec	~\$0.18/search	Day-to-day use (recommended)
Fast	Basic	~5 sec	~\$0.05/search	Quick exploration, brainstorming

Select your mode before running a search. All modes use the same methodology—the difference is in the AI model powering each stage.

Cost Transparency

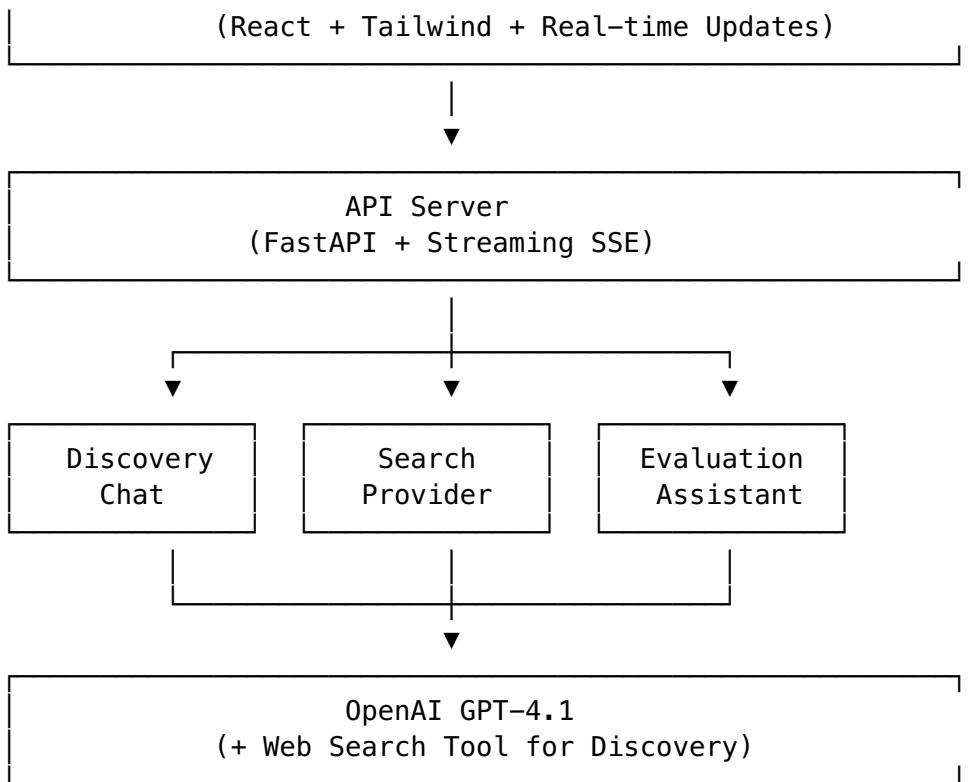
Every operation shows real-time cost tracking:

- Discovery chat: ~\$0.02/message
- Search (50 candidates): ~\$0.80-1.50
- Evaluation: ~\$0.15-0.30
- Typical full session: **\$1-2 total**

No hidden fees. No subscriptions. Pay only for what you use.

Architecture

Web Interface



Key Design Decisions:

- Single LLM with specialized prompts (not multi-agent)
- Batch processing to avoid quality degradation at scale
- External state management for consistent scoring
- Human-in-the-loop at every stage

What Makes PartnerScope Different

Traditional Approach	PartnerScope
Manual database searches	AI-powered multi-source discovery
Spreadsheet tracking	Structured evaluation with rankings
Gut-feel decisions	Multi-dimensional scoring with confidence
One-time analysis	Dynamic refinement without re-running
No validation	Comparison against external research tools

Use Cases

Pilot Partners > “We need universities or employers who can provide 100-300 users for an 8-12 week pilot”

Distribution Partners > “Looking for healthcare distributors with hospital network relationships”

Technology Partners > “Need API integration partners in the fintech space”

Getting Started

1. **Start a conversation** — Tell us about your startup
2. **Review your profile** — Edit the extracted requirements
3. **Run search** — Select data sources and quality level
4. **Evaluate results** — Customize dimensions and weights
5. **Compare & refine** — Validate against external tools, adjust as needed

The entire process takes 5-15 minutes for a comprehensive partner search that would traditionally require days of manual research.

Technical Specifications

- **Backend:** Python, FastAPI
 - **Frontend:** React, Tailwind CSS
 - **AI Models:** OpenAI GPT-4.1 series (configurable)
 - **Data Sources:** CSV/Database import, Real-time web search
 - **Output Formats:** JSON, exportable rankings
-

Research Foundation

PartnerScope draws inspiration from academic research in partner selection:

- **PartnerMAS** (arXiv:2509.24046): Multi-agent framework for partner evaluation
- **Batch Processing**: Inspired by insights from Recursive Language Models (arXiv:2512.24601)

We've adapted these approaches for practical startup use—prioritizing interactivity and human control over full automation.

PartnerScope — From months of research to minutes of conversation.

UTokyo Research, 2026